

Data Science R Basics

Learning Objectives

- Learn to read, extract, and create datasets in R
- Learn to perform a variety of operations on datasets using R
- Learn to write your own functions/sub-routines in R

Course Overview

Section 1: R Basics, Functions, Data types

You will get started with R, learn about its functions and data types.

Section 2: Vectors, Sorting

You will learn to operate on vectors and advanced functions such as sorting.

Section 3: Indexing, Data Manipulation, Plots

You will learn to wrangle and visualize data.

Section 4: Programming Basics

You will learn to use general programming features like ‘if-else’, and ‘for loop’ commands, and write your own functions to perform various operations on datasets.

Section 1 Overview

Section 1 introduces you to R Basics, Functions and Datatypes.

In Section 1, you will learn to:

- Appreciate the rationale for data analysis using R
- Define objects and perform basic arithmetic and logical operations
- Use pre-defined functions to perform operations on objects
- Distinguish between various data types

The textbook for this section is available [here](#)

Motivation

Here is a link to the textbook section on the motivation for this course.

Getting started

Here is a link to the textbook section on Getting Started with R.

Key Points

- R was developed by statisticians and data analysts as an interactive environment for data analysis.
- Some of the advantages of R are that (1) it is free and open source, (2) it has the capability to save scripts, (3) there are numerous resources for learning, and (4) it is easy for developers to share software implementation.
- Expressions are evaluated in the R console when you type the expression into the console and hit Return.
- A great advantage of R over point and click analysis software is that you can save your work as scripts.
- “Base R” is what you get after you first install R. Additional components are available via packages.

```
# installing the dslabs package  
if(!require(dslabs)) install.packages("dslabs")
```

```
## Loading required package: dslabs
```

```
# loading the dslabs package into the R session  
library(dslabs)
```

Installing R and R Studio

Installing R To install R to work on your own computer, you can download it freely from the Comprehensive R Archive Network (CRAN). Note that CRAN makes several versions of R available: versions for multiple operating systems and releases older than the current one. You want to read the CRAN instructions to assure you download the correct version. If you need further help, you read the walkthrough in this Chapter of the textbook.

Installing RStudio RStudio is an integrated development environment (IDE). We highly recommend installing and using RStudio to edit and test your code. You can install RStudio through the RStudio website. Their cheatsheet is a great resource. You must install R before installing RStudio.

Textbook Link Here is a link to the textbook section on Installing R and RStudio.

R Basics

Objects

Here is a link to the textbook section on objects in R.

Key Points

- To define a variable, we may use the assignment symbol “<-”.
- There are two ways to see the value stored in a variable: (1) type the variable into the console and hit Return, or (2) type print(“variable name”) and hit Return.
- Objects are stuff that is stored in R. They can be variables, functions, etc.
- The ls() function shows the names of the objects saved in your workspace.

Solving the equation $x^2+x-1=0$

```
# assigning values to variables
a <- 1
b <- 1
c <- -1

# solving the quadratic equation
(-b + sqrt(b^2 - 4*a*c) ) / ( 2*a )
```

```
## [1] 0.618034
```

```
(-b - sqrt(b^2 - 4*a*c) ) / ( 2*a )
```

```
## [1] -1.618034
```

Functions

Here is a link to the textbook section on functions.

Key points

- In general, to evaluate a function we need to use parentheses. If we type a function without parentheses, R shows us the code for the function. Most functions also require an argument, that is, something to be written inside the parenthesis.
- To access help files, we may use the help function `help("function name")`, or write the question mark followed by the function name.
- The help file shows you the arguments the function is expecting, some of which are required and some are optional. If an argument is optional, a default value is assigned with the equal sign. The `args()` function also shows the arguments a function needs.
- To specify arguments, we use the equals sign. If no argument name is used, R assumes you're entering arguments in the order shown in the help file.
- Creating and saving a script makes code much easier to execute.
- To make your code more readable, use intuitive variable names and include comments (using the “#” symbol) to remind yourself why you wrote a particular line of code.

Assessment - R Basics

1. What is the sum of the first n positive integers? We can use the formula $n(n+1)/2$ to quickly compute this quantity.

```
# Here is how you compute the sum for the first 20 integers
20*(20+1)/2
```

```
## [1] 210
```

```
# However, we can define a variable to use the formula for other values of n
n <- 20
n*(n+1)/2
```

```
## [1] 210
```

```
n <- 25
n*(n+1)/2
```

```
## [1] 325
```

```
# Below, write code to calculate the sum of the first 100 integers
n<-100
n*(n+1)/2
```

```
## [1] 5050
```

2. What is the sum of the first 1000 positive integers? We can use the formula $n(n+1)/2$ to quickly compute this quantity.

```
# Below, write code to calculate the sum of the first 1000 integers
n<-1000
n*(n+1)/2
```

```
## [1] 500500
```

3. Run the following code in the R console.

```
n <- 1000
x <- seq(1, n)
sum(x)
```

```
## [1] 500500
```

Based on the result, what do you think the functions `seq` and `sum` do?

- ☐ A. `sum` creates a list of numbers and `seq` adds them up.
- ☒ B. `seq` creates a list of numbers and `sum` adds them up.
- ☐ C. `seq` computes the difference between two arguments and `sum` computes the sum of 1 through 1000.
- ☐ D. `sum` always returns the same number.

4. In math and programming we say we evaluate a function when we replace arguments with specific values. So if we type `log2(16)` we evaluate the `log2` function to get the log base 2 of 16 which is 4.

In R it is often useful to evaluate a function inside another function. For example, `sqrt(log2(16))` will calculate the log to the base 2 of 16 and then compute the square root of that value. So the first evaluation gives a 4 and this gets evaluated by `sqrt` to give the final answer of 2.

```
# log to the base 2
log2(16)
```

```
## [1] 4
```

```
# sqrt of the log to the base 2 of 16:
sqrt(log2(16))
```

```
## [1] 2
```

```
# Compute log to the base 10 (log10) of the sqrt of 100. Do not use variables.
log10(sqrt(100))
```

```
## [1] 1
```

5. Which of the following will always return the numeric value stored in `x`? You can try out examples and use the help system in the R console.

- ☐ A. `log(10^x)`
- ☐ B. `log10(x^10)`
- ☒ C. `log(exp(x))`
- ☐ D. `exp(log(x, base = 2))`

Assessment 2

1. Load the US murders dataset.

```
library(dslabs)
data(murders)
```

Use the function `str` to examine the structure of the `murders` object. We can see that this object is a data frame with 51 rows and five columns. Which of the following best describes the variables represented in this data frame? - [] A. The 51 states. - [] B. The murder rates for all 50 states and DC. - [x] C. The state name, the abbreviation of the state name, the state's region, and the state's population and total number of murders for 2010. - [] D. `str` shows no relevant information.

2. What are the column names used by the data frame for these five variables?

```
# Load package and data
library(dslabs)
data(murders)
```

```
# Use the function names to extract the variable names
names(murders)
```

```
## [1] "state"      "abb"        "region"     "population" "total"
```

3. Use the accessor `$` to extract the state abbreviations and assign them to the object `a`. What is the class of this object?

```
# To access the population variable from the murders dataset use this code:
p <- murders$population
```

```
# To determine the class of object `p` we use this code:
class(p)
```

```
## [1] "numeric"
```

```
# Use the accessor to extract state abbreviations and assign it to a
a <- murders$abb
```

```
# Determine the class of a
class(a)
```

```
## [1] "character"
```

4. Now use the square brackets to extract the state abbreviations and assign them to the object b. Use the identical function to determine if a and b are the same.

```
# We extract the population like this:
p <- murders$population
```

```
# This is how we do the same with the square brackets:
o <- murders[["population"]]
```

```
# We can confirm these two are the same
identical(o, p)
```

```
## [1] TRUE
```

```
# Use square brackets to extract `abb` from `murders` and assign it to b
b<-murders[["abb"]]
```

```
# Check if `a` and `b` are identical
identical(a, b)
```

```
## [1] TRUE
```

5. We saw that the region column stores a factor. You can corroborate this by typing:

```
class(murders$region)
```

With one line of code, use the function levels and length to determine the number of regions defined by this dataset.

```
# We can see the class of the region variable using class
class(murders$region)
```

```
## [1] "factor"
```

```
# Determine the number of regions included in this variable
length(levels(murders$region))
```

```
## [1] 4
```

6. The function `table` takes a vector and returns the frequency of each element. You can quickly see how many states are in each region by applying this function. Use this function in one line of code to create a table of states per region.

```
# Here is an example of what the table function does
x <- c("a", "a", "b", "b", "b", "c")
table(x)

## x
## a b c
## 2 3 1

# Write one line of code to show the number of states per region
table(murders$region)

##
##      Northeast      South North Central      West
##           9           17           12           13
```

Section 2 Overview

In Section 2.1, you will: - Create numeric and character vectors. - Name the columns of a vector. - Generate numeric sequences. - Access specific elements or parts of a vector. - Coerce data into different data types as needed.

In Section 2.2, you will: - Sort vectors in ascending and descending order. - Extract the indices of the sorted elements from the original vector. - Find the maximum and minimum elements, as well as their indices, in a vector. - Rank the elements of a vector in increasing order.

In Section 2.3, you will: - Perform arithmetic between a vector and a single number. - Perform arithmetic between two vectors of same length.

The textbook for this section is available [here](#)

Assessment 3

1. Use the function `c` to create a vector with the average high temperatures in January for Beijing, Lagos, Paris, Rio de Janeiro, San Juan and Toronto, which are 35, 88, 42, 84, 81, and 30 degrees Fahrenheit. Call the object `temp`.

```
# Here is an example creating a numeric vector named cost
cost <- c(50, 75, 90, 100, 150)
```

```
# Create a numeric vector to store the temperatures listed in the instructions into a vector named temp
# Make sure to follow the same order in the instructions
temp <- c(35, 88, 42, 84, 81, 30)
```

2. Now create a vector with the city names and call the object `city`.

```
# here is an example of how to create a character vector
food <- c("pizza", "burgers", "salads", "cheese", "pasta")
```

```
# Create a character vector called city to store the city names
# Make sure to follow the same order as in the instructions
city <- c("Beijing", "Lagos", "Paris", "Rio de Janeiro", "San Juan", "Toronto")
```

3. Use the names function and the objects defined in the previous exercises to associate the temperature data with its corresponding city.

```
# Associate the cost values with its corresponding food item
cost <- c(50, 75, 90, 100, 150)
food <- c("pizza", "burgers", "salads", "cheese", "pasta")
names(cost) <- food

# You already wrote this code
temp <- c(35, 88, 42, 84, 81, 30)
city <- c("Beijing", "Lagos", "Paris", "Rio de Janeiro", "San Juan", "Toronto")

# Associate the temperature values with its corresponding city
names(temp) <- city
```

4. Use the [and : operators to access the temperature of the first three cities on the list.

```
# cost of the last 3 items in our food list:
cost[3:5]

## salads cheese  pasta
##      90      100      150

# temperatures of the first three cities in the list:
temp[1:3]

## Beijing  Lagos  Paris
##      35      88      42
```

5. Use the [operator to access the temperature of Paris and San Juan.

```
# Access the cost of pizza and pasta from our food list
cost[c(1,5)]

## pizza pasta
##      50      150

# Define temp
temp <- c(35, 88, 42, 84, 81, 30)
city <- c("Beijing", "Lagos", "Paris", "Rio de Janeiro", "San Juan", "Toronto")
names(temp) <- city

# Access the temperatures of Paris and San Juan
temp[c(3,5)]

##      Paris San Juan
##      42      81
```

6. Use the : operator to create a sequence of numbers 12, 13, 14,...,73.


```
# Create a vector m of integers that starts at 32 and ends at 99.
m <- 32:99
```

```
# Determine the length of object m.
length(m)
```

```
## [1] 68
```

```
# Create a vector x of integers that starts 12 and ends at 73.
x <- 12:73
```

```
# Determine the length of object x.
length(x)
```

```
## [1] 62
```

7. Create a vector containing all the positive odd numbers smaller than 100.

```
# Create a vector with the multiples of 7, smaller than 50.
seq(7, 49, 7)
```

```
## [1] 7 14 21 28 35 42 49
```

```
# Create a vector containing all the positive odd numbers smaller than 100.
# The numbers should be in ascending order
seq(1,100,2)
```

```
## [1] 1 3 5 7 9 11 13 15 17 19 21 23 25 27 29 31 33 35 37 39 41 43 45
## [24] 47 49 51 53 55 57 59 61 63 65 67 69 71 73 75 77 79 81 83 85 87 89 91
## [47] 93 95 97 99
```

8. Create a vector of numbers that starts at 6, does not pass 55, and adds numbers in increments of $4/7$: 6, $6+4/7$, $6+8/7$, etc.. How many numbers does the list have? Hint: use seq and length.

```
# We can a vector with the multiples of 7, smaller than 50 like this
seq(7, 49, 7)
```

```
## [1] 7 14 21 28 35 42 49
```

```
# But note that the second argument does not need to be last number.
# It simply determines the maximum value permitted.
# so the following line of code produces the same vector as seq(7, 49, 7)
seq(7, 50, 7)
```

```
## [1] 7 14 21 28 35 42 49
```

```
# Create a sequence of numbers from 6 to 55, with  $4/7$  increments and determine its length
length(seq(6,55,4/7))
```

```
## [1] 86
```

9. What is the class of the following object `a <- seq(1, 10, length.out = 100)`?

```
# Store the sequence in the object a
a <- seq(1, 10, length.out = 100)

# Determine the class of a
class(a)

## [1] "numeric"
```

10. What is the class of the following object `a <- seq(1, 10)`?

```
# Store the sequence in the object a
a <- seq(1, 10)

# Determine the class of a
class(a)

## [1] "integer"
```

11. The class of `class(a<-1)` is numeric, not integer. R defaults to numeric and to force an integer, you need to add the letter L. Confirm that the class of `1L` is integer.

```
# Check the class of 1, assigned to the object a
a <- class(1)

# Confirm the class of 1L is integer
class(1L)

## [1] "integer"
```

12. Define the following vector:

```
x <- c("1", "3", "5", "a")
```

and coerce it to get integers.

```
# Define the vector x
x <- c(1, 3, 5,"a")

# Note that the x is character vector
x

## [1] "1" "3" "5" "a"

# Typecast the vector to get an integer vector
x <- as.numeric(x)

# You will get a warning but that is ok
```

Assessment 4

For these exercises we will use the US murders dataset. Make sure you load it prior to starting.

```
library(dslabs)
data("murders")
```

1. Use the \$ operator to access the population size data and store it as the object pop. Then use the sort function to redefine pop so that it is sorted. Finally, use the [operator to report the smallest population size.

```
# Access the `state` variable and store it in an object
states <- murders$state

# Sort the object alphabetically and redefine the object
states <- sort(states)

# Report the first alphabetical value
states[1]

## [1] "Alabama"

# Access population values from the dataset and store it in pop
pop <- murders$population

# Sort the object and save it in the same object
pop<-sort(pop)

# Report the smallest population size
pop[1]

## [1] 563626
```

2. Now instead of the smallest population size, find the index of the entry with the smallest population size. Hint: use order instead of sort.

```
# Access population from the dataset and store it in pop
pop <- murders$population

# Use the command order, to order pop and store in object o
o <- order(pop)

# Find the index number of the entry with the smallest population size
which.min(murders$population)

## [1] 51
```

3. We can actually perform the same operation as in the previous exercise using the function which.min. Write one line of code that does this.

```
# Find the smallest value for variable total
which.min(murders$total)
```

```
## [1] 46
```

```
# Find the smallest value for population
which.min(murders$population)
```

```
## [1] 51
```

4. Now we know how small the smallest state is and we know which row represents it. Which state is it? Define a variable `states` to be the state names from the `murders` data frame. Report the name of the state with the smallest population.

```
# Define the variable i to be the index of the smallest state
i <- which.min(murders$population)
```

```
# Define variable states to hold the states
states <- murders$state
```

```
# Use the index you just defined to find the state with the smallest population
states[i]
```

```
## [1] "Wyoming"
```

5. You can create a data frame using the `data.frame` function. Here is a quick example:

```
temp <- c(35, 88, 42, 84, 81, 30)
city <- c("Beijing", "Lagos", "Paris", "Rio de Janeiro", "San Juan", "Toronto")
city_temps <- data.frame(name = city, temperature = temp)
```

Use the `rank` function to determine the population rank of each state from smallest population size to biggest. Save these ranks in an object called `ranks`, then create a data frame with the state name and its rank. Call the data frame `my_df`.

```
# Store temperatures in an object
temp <- c(35, 88, 42, 84, 81, 30)
```

```
# Store city names in an object
city <- c("Beijing", "Lagos", "Paris", "Rio de Janeiro", "San Juan", "Toronto")
```

```
# Create data frame with city names and temperature
city_temps <- data.frame(name = city, temperature = temp)
```

```
# Define a variable states to be the state names
states <- murders$state
```

```
# Define a variable ranks to determine the population size ranks
ranks <- rank(murders$population)
```

```
# Create a data frame my_df with the state name and its rank
my_df <- data.frame(name=states, ranks)
```

6. Repeat the previous exercise, but this time order `my_df` so that the states are ordered from least populous to most populous. Hint: create an object `ind` that stores the indexes needed to order the population values. Then use the bracket operator `[` to re-order each column in the data frame.

```

# Define a variable states to be the state names from the murders data frame
states <- murders$state

# Define a variable ranks to determine the population size ranks
ranks <- rank(murders$population)

# Define a variable ind to store the indexes needed to order the population values
ind <- order(murders$population)

# Create a data frame my_df with the state name and its rank and ordered from least populous to most
my_df<-data.frame(states = states[ind], ranks = ranks[ind])

```

7. The `na_example` vector represents a series of counts. You can quickly examine the object using:

```

data("na_example")
str(na_example)
nt [1:1000] 2 1 3 2 1 3 1 4 3 2 ...

```

However, when we compute the average with the function `mean`, we obtain an NA:

```

mean(na_example)
[1] NA

```

The `is.na` function returns a logical vector that tells us which entries are NA. Assign this logical vector to an object called `ind` and determine how many NAs does `na_example` have.

```

# Using new dataset
library(dslabs)
data(na_example)

# Checking the structure
str(na_example)

## int [1:1000] 2 1 3 2 1 3 1 4 3 2 ...

# Find out the mean of the entire dataset
mean(na_example)

## [1] NA

# Use is.na to create a logical index ind that tells which entries are NA
ind <- is.na(na_example)

# Determine how many NA ind has using the sum function
sum(ind)

## [1] 145

```

8. Now compute the average again, but only for the entries that are not NA. Hint: remember the `!` operator.

```

# Note what we can do with the ! operator
x <- c(1, 2, 3)
ind <- c(FALSE, TRUE, FALSE)
x[!ind]

## [1] 1 3

# Create the ind vector
library(dslabs)

data(na_example)
ind <- is.na(na_example)

# We saw that this gives an NA
mean(na_example)

## [1] NA

# Compute the average, for entries of na_example that are not NA
mean(na_example[!ind])

## [1] 2.301754

```

Assessment 5

1. Previously we created this data frame:

```

temp <- c(35, 88, 42, 84, 81, 30)
city <- c("Beijing", "Lagos", "Paris", "Rio de Janeiro", "San Juan", "Toronto")
city_temps <- data.frame(name = city, temperature = temp)

```

Remake the data frame using the code above, but add a line that converts the temperature from Fahrenheit to Celsius. The conversion is $C = 5/9 * (F - 32)$.

```

# Assign city names to `city`
city <- c("Beijing", "Lagos", "Paris", "Rio de Janeiro", "San Juan", "Toronto")

# Store temperature values in `temp`
temp <- c(35, 88, 42, 84, 81, 30)

# Convert temperature into Celsius and overwrite the original values of 'temp' with these Celsius values
temp <- 5/9*(temp-32)

# Create a data frame `city_temps`
city_temps <- data.frame(name=city,temperature=temp)

```

2. What is the following sum $1 + 1/2^2 + 1/3^2 + .1/100^2$? Hint: thanks to Euler, we know it should be close to $??^2/6$.

```
# Define an object `x` with the numbers 1 through 100
x <- c(1:100)
```

```
# Compute the sum
sum(1/x^2)
```

```
## [1] 1.634984
```

3. Compute the per 100,000 murder rate for each state and store it in the object `murder_rate`. Then compute the average murder rate for the US using the function `mean`. What is the average?

```
# Load the data
library(dslabs)
data(murders)
```

```
# Store the per 100,000 murder rate for each state in murder_rate
murder_rate <- murders$total/murders$population *100000
```

```
# Calculate the average murder rate in the US
mean(murder_rate)
```

```
## [1] 2.779125
```

Section 3 Overview

Section 3 introduces to the R commands and techniques that help you wrangle, analyze, and visualize data.

In Section 3.1, you will:

- Subset a vector based on properties of another vector.
- Use multiple logical operators to index vectors.
- Extract the indices of vector elements satisfying one or more logical conditions.
- Extract the indices of vector elements matching with another vector.
- Determine which elements in one vector are present in another vector.

In Section 3.2, you will:

- Wrangle data tables using the functions in ‘dplyr’ package.
- Modify a data table by adding or changing columns.
- Subset rows in a data table.
- Subset columns in a data table.
- Perform a series of operations using the pipe operator.
- Create data frames.

In Section 3.3, you will:

- Plot data in scatter plots, box plots and histograms.

The textbook for this section is available [here](#)

Assessment 6

Start by loading the library and data.

```
library(dslabs)
data(murders)
```

1. Compute the per 100,000 murder rate for each state and store it in an object called `murder_rate`. Then use logical operators to create a logical vector named `low` that tells us which entries of `murder_rate` are lower than 1.

```
# Store the murder rate per 100,000 for each state, in `murder_rate`
murder_rate <- murders$total / murders$population * 100000
```

```
# Store the `murder_rate < 1` in `low`
low <- murder_rate < 1
```

2. Now use the results from the previous exercise and the function which to determine the indices of murder_rate associated with values lower than 1.

```
# Store the murder rate per 100,000 for each state, in murder_rate
murder_rate <- murders$total/murders$population*100000
```

```
# Store the murder_rate < 1 in low
low <- murder_rate < 1
```

```
# Get the indices of entries that are below 1
which(low)
```

```
## [1] 12 13 16 20 24 30 35 38 42 45 46 51
```

3. Use the results from the previous exercise to report the names of the states with murder rates lower than 1.

```
# Store the murder rate per 100,000 for each state, in murder_rate
murder_rate <- murders$total/murders$population*100000
```

```
# Store the murder_rate < 1 in low
low <- murder_rate < 1
```

```
# Names of states with murder rates lower than 1
murders$state[low]
```

```
## [1] "Hawaii"      "Idaho"       "Iowa"        "Maine"
## [5] "Minnesota"   "New Hampshire" "North Dakota" "Oregon"
## [9] "South Dakota" "Utah"        "Vermont"     "Wyoming"
```

4. Now extend the code from exercise 2 and 3 to report the states in the Northeast with murder rates lower than 1. Hint: use the previously defined logical vector low and the logical operator &.

```
# Store the murder rate per 100,000 for each state, in `murder_rate`
murder_rate <- murders$total/murders$population*100000
```

```
# Store the `murder_rate < 1` in `low`
low <- murder_rate < 1
```

```
# Create a vector ind for states in the Northeast and with murder rates lower than
ind <- low & murders$region=='Northeast'
```

```
# Names of states in `ind`
murders$state[ind]
```

```
## [1] "Maine"      "New Hampshire" "Vermont"
```


5. In a previous exercise we computed the murder rate for each state and the average of these numbers. How many states are below the average?

```
# Store the murder rate per 100,000 for each state, in murder_rate
murder_rate <- murders$total/murders$population*100000

# Compute average murder rate and store in avg using `mean`
avg <- mean(murder_rate)

# How many states have murder rates below avg ? Check using sum
sum(murder_rate<avg)

## [1] 27
```

6. Use the match function to identify the states with abbreviations AK, MI, and IA. Hint: start by defining an index of the entries of murders\$abb that match the three abbreviations, then use the [operator to extract the states.

```
# Store the 3 abbreviations in abbs in a vector (remember that they are character vectors and need quotes)
abbs <- c('AK','MI','IA')

# Match the abbs to the murders$abb and store in ind
ind <- match(abbs , murders$abb)

# Print state names from ind
murders$state[ind]

## [1] "Alaska" "Michigan" "Iowa"
```

7. Use the %in% operator to create a logical vector that answers the question: which of the following are actual abbreviations: MA, ME, MI, MO, MU ?

```
# Store the 5 abbreviations in `abbs`. (remember that they are character vectors)
abbs <- c('MA', 'ME', 'MI', 'MO', 'MU')

# Use the %in% command to check if the entries of abbs are abbreviations in the the murders data frame
abbs%in%murders$abb

## [1] TRUE TRUE TRUE TRUE FALSE
```

8. Extend the code you used in exercise 7 to report the one entry that is not an actual abbreviation. Hint: use the ! operator, which turns FALSE into TRUE and vice versa, then which to obtain an index.

```
# Store the 5 abbreviations in abbs. (remember that they are character vectors)
abbs <- c("MA", "ME", "MI", "MO", "MU")

# Use the `which` command and `!` operator to find out which abbreviation are not actually part of the
ind <- which(!abbs%in%murders$abb)

# What are the entries of abbs that are not actual abbreviations
abbs[ind]

## [1] "MU"
```